

National University of Singapore
School of Computing

CS2105

Assignment 1 (8 Marks)

Semester 1 AY21/22

Submission Deadline

24th Sept 2021 (Friday) 11:59pm sharp. 2 marks penalty will be imposed on late submission (Late submission refers to submission or re-submission after the deadline). The submission folder will be closed on **1st Oct 2021 (Friday) 11:59pm** sharp and no late submission will be accepted afterwards.

Objectives

In this assignment, you will implement a server that communicates with clients using an HTTP-like protocol. After completing this assignment, you should

- be able to implement a simple TCP-based server, and
- have a good understanding of the mechanisms of HTTP.

Group Work

All the work in this assignment should be done individually. However, if you find the assignment too difficult,

- you are allowed to form a group with another student
- **maximum two students per group.**
- Group submission is subject to **2 marks penalty** for each student.

Under no circumstances should you solve it in a group and then submit it as an individual solution. This is considered plagiarism. **There will be no acceptance for excuses such as forgetting to declare as group submission.** Please refer to the "Special Instructions for Group Submission" and "Plagiarism Warning" on page 4 for more details.

Grading

We will test and grade your program on the sunfire server. Please make sure that your program run properly on sunfire. Moreover, you are allowed to use libraries installed in public folders of sunfire (e.g. /usr/lib) only.

We accept submission of Python 3 (**in particular, 3.7**), Java, or C/C++ program, and we recommend that you use **Python 3** for your assignments. Programming languages other than Python 3, Java, and C/C++ are not allowed. For Python 3, we use the python3 program installed in folder /usr/local/Python-3.7/bin on sunfire for grad-

ing. If you use Java or C/C++, we will compile and run your program for grading using the default compilers on sunfire (java 9.0.4 installed in /usr/local/java/jdk/bin, or gcc 4.8.4 installed in /usr/local/gcc-4.8/bin). The grading script infers your programming language from the file extension name (.py, .java, .c). Therefore, please ensure your files have the correct extension names. For a Java program, the class name should be consistent with the source file name, and please implement the static `main()` method so that the program can be executed as a standalone process after compilation. We will **deduct 1 mark** for every type of failure to follow instructions (e.g. wrong program name).

Note that for **Java** programs, name your main class as `WebServer` and hence source file name `WebServer.java` during development. When you want to test your program using the provided script on Sunfire, or to make submission, rename your file name according to program submission and group submission section **while keeping the class name as `WebServer`**. Grading script will rename your file accordingly during compilation.

We will grade your program based on its correctness only. A grading script will be used to test your program and no manual grading will be provided.

Testing Your Program

To test your program, please use your SoC UNIX ID and password to log on to sunfire as instructed on Assignment 0 paper. To make the `python3` alias permanently and thus avoid typing the `alias` command every time you login, you can run the following command **once** to store the shortcut into the configuration file:

```
echo alias python3=/usr/local/Python-3.7/bin/python3 >> ~/.bash_profile
```

- Your server program should receive one **command-line argument** which is the port number to `listen()`, as the following command shows:

```
python3 WebServer-A0165432X.py <port>
```

- If you test your server manually, please select a port number greater than 1024, because the OS usually restricts usage of ports lower than 1024. If you get a `BindException: Address already in use` (or similar errors for other languages), please try a different port number. Note that your program should not read from `stdin`. Your program can print anything to `stdout` or `stderr`, and our test script will silently ignore them.
- We also release a set of grading scripts to you under the **test** folder. **There is no hidden test cases during grading**. However, passing all the test cases does not guarantee that you will get full marks. We will detect fraudulent cases such as hard-coding answers.
- To use the grading script, please upload your program along with the `test` folder given in the package to sunfire. Make sure that your program and the `test` folder are in the same directory. Then, you can run the following command to test your

server program:

```
bash test/WebServer.sh
```

- By default, the script runs through all test cases. You can also choose to run a certain test case by specifying the case number in the command:

```
bash test/WebServer.sh 3
```

To stop a test, press `ctrl-c`. If pressing the key combination once does not work, hold the keys until the script exits.

- If you ever encounter this error: **tput: unknown terminal "xterm-256color"** when testing your program using script provided, run the command:

```
export TERM=xterm
```

once after you log in and before you run WebServer.sh.

There is a low possibility that the grading script is grading an existing running server program instead of yours, if they use same port. So start the assignment early to avoid "congestion" during last few days. An advice to ensure the grading feedback you received is on your current program is to print to screen your student number right after receiving a connection, and kill your own existing server programs if any before testing it again. The grading result will be based on your program if you see your student number printed for particular test case, or otherwise try again.

Program Submission

For individual submission, please name your single source file as `WebServer-<Matric number>.py` and submit it to the `Assignment_1_student_submission` folder of LumiNUS Files. Here, `<Matric number>` is your matriculation number which starts with letter A. An example file name would be `WebServer-A0165432X.py`. If you use Java, C, or C++ to implement the web server, please use `.java`, `.c`, or `.cpp` respectively as the extension name. Note that file names are **case-sensitive** on sunfire.

You are not allowed to post your solutions to any publicly accessible site on the Internet.

Special Instructions for Group Submission

For group submission, please include matriculation numbers of both students in the file name, i.e. `WebServer-<Matric number 1>-<Matric number 2>.py`. Submit it to the same `Assignment_1_student_submission` folder. An example file name would be `WebServer-A0165432X-A0123456Y.py`. For each group, there should be one designated member who submits the file, to avoid problems caused by multiple branches within a group. **Do not change the designated submitter!** If the group needs to upload a new version, it should be done by the same designated submitter as well.

Plagiarism Warning

You are free to discuss this assignment with your friends. However, you should refrain from sharing your program, program fragments, or detailed algorithms with others. If you want to solve this assignment in a group, please do so and **declare it as group work**.

We employ zero-tolerance policy against plagiarism. If a suspicious case is found, student would be asked to explain his/her code to the evaluator in face. Confirmed breach may result in zero mark for the assignment and further disciplinary action from the school.

Question & Answer

If you have any doubts on this assignment, please post your questions on LumiNUS forum before consulting the teaching team. However, the teaching team will NOT debug programs for students and we provide support for language-specific questions as a best-effort service. The intention of Q&A is to help clarify misconceptions or give you necessary directions.

The Server

In this assignment, you are required to implement an HTTP-like server that provides key-value store services and a frequency counter service. The server consists of two components. The first is a custom protocol simplified from HTTP 1.1. It retains the basic request-response mechanism and the persistent connection feature of HTTP 1.1, but uses a different header format. On top of this protocol, an interactive in-memory key-value store and a frequency counter are to be implemented under the path `/key/` and `/counter/` respectively. Figure 1 shows the overall architecture of the web server.

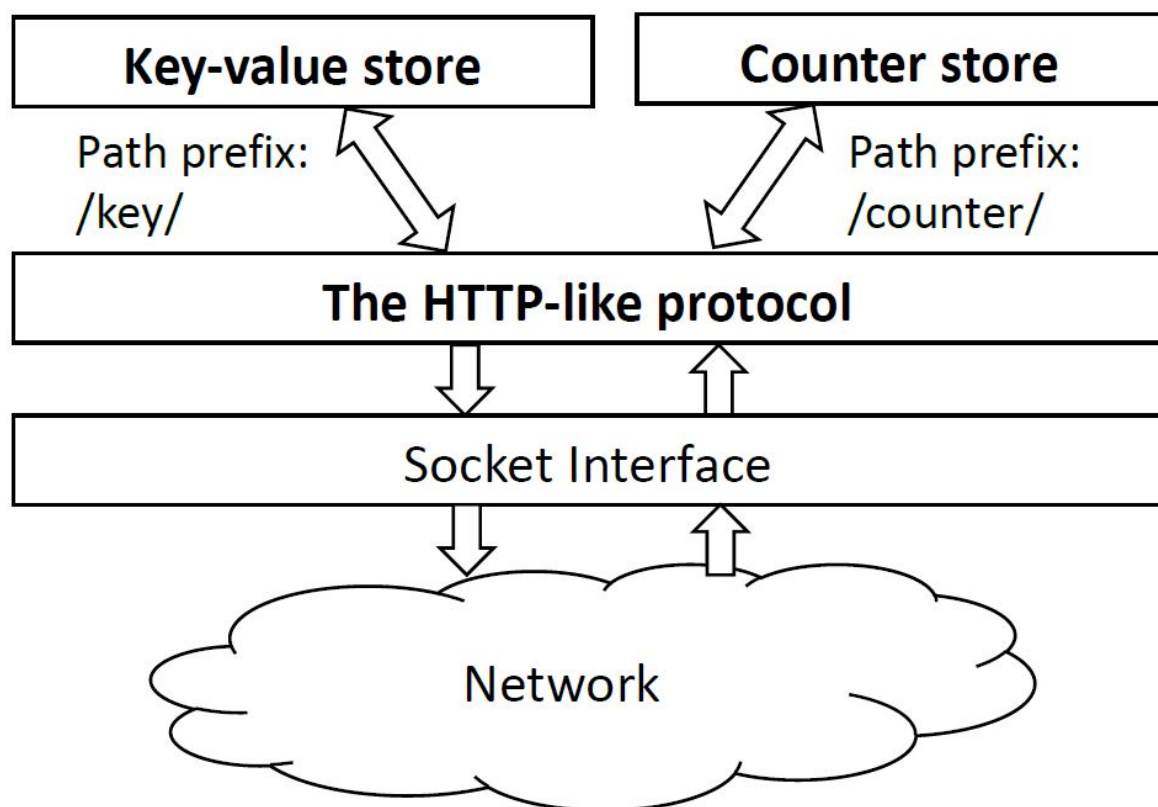


Figure 1: The layered architecture of the server. **You need to implement the top two layers which are highlighted with bold fonts.**

The TCP Socket

The lowest-level layer of this server is the TCP socket interface. Basically, the server should

1. `bind()` to a port
2. `listen()` to the socket and `accept()` a new connection
3. Parse and respond to client request(s) sequentially, by using `recv()` and `sendall()` for socket I/O.
4. Going back to 2 after the client disconnects

As discussed in the "Testing Your Program" section, port number is passed to your server as **a command line argument**. The grading script ensures that only one client connects to your server at any time so that no connection multiplexing or polling is required; your server can simply process connections sequentially.

It is necessary to detect disconnection events reliably, because your server is required to support the persistent connection feature originated from HTTP 1.1. If the `bytes` object returned by `recv()` is of zero length, then no more data could be `recv()`'ed from the connection. Since this only indicates half close (from client side only) of a TCP connection, your server should finish processing the final request(s) before closing the connection on server side. **The client is guaranteed to keep the connection half open until either the server finishes sending all responses or a timeout occurs. Therefore, there will be no `sendall()` failures on the server side, unless it times out during this process.**

The HTTP-like Protocol

After establishing a new TCP connection with a client, your server should read a customized HTTP request from the socket, which consists of request headers and an optional request body. Your server parses the request and sends the corresponding response as defined below. After such a request-response cycle, the connection is **persisted** (i.e. not closing immediately), and the server waits for another request from the same client until a notification of socket disconnection.

Since a real-world TCP connection could be intermittent, your server is also required to work properly when requests are delivered in chunks of random sizes with delays between chunks. Note that a TCP connection only guarantees ordering of data, but not segment sizes or delays. The test script simulates this type of intermittent connection in some test cases, while keeping the delays between two consecutive transfers below one second. To handle chunks that contain incomplete requests, it is sufficient for your server to handle data buffering well.

In addition, your server should be **responsive** in the sense that upon receiving a complete request, the server processes it and sends back the response immediately. The test script sets a one-second timeout after each request is **fully** sent to the server. If a single chunk contains multiple requests, your server is also expected to respond quickly to all complete requests inside.

To grasp the core ideas of HTTP, we omit other tedious aspects of the protocol. For this assignment, the server only needs to handle basic request and response information, as well as the `Content-Length` header field. The **header format** of this HTTP-like protocol is simplified as follows:

1. A header is a string consisting of **non-empty substrings** delimited with white-spaces (ASCII code 32). Two consecutive white-spaces mark the end of a header. You may assume that a substring does not contain white-spaces. Obviously, our header format is different from (and simpler than!) that of standard HTTP (e.g.,

no more \r\n or colons).

2. You may assume that a header consists of at least two substrings. The first two substrings (compulsory) contain information specific to requests or responses. Additional (and optional) substrings may follow. Every two substrings form a header field, where the first substring is the **case-insensitive** name of the header field, and the second is the value of this field.
3. For a request, the first two substrings (compulsory) are *HTTP method* and *path*, respectively. The HTTP method is **case-insensitive**, while path is **case-sensitive**. After these two substrings, optional header fields may follow, e.g. Content-Length(**case-insensitive**) and its value (a non-negative integer) if there is a content body in the request. The header finally ends with two white-spaces. This is followed by the content body, if any. Some sample requests are shown below:

- (a) GET_/key/CS2105_
- (b) POST_/key/ModuleCode_Content-Length_6_ CS2105
- (c) GET_/key/ModuleCode_
- (d) GET_/counter/CS2105_
- (e) POST_/counter/StudentNumber_
- (f) GET_/counter/StudentNumber_
- (g) DELETE_/key/ModuleCode_
- (h) DELETE_/counter/CS2105_

(Here, space is illustrated by _ and content body is underlined.)

4. For a response, the first two substrings (compulsory) are *HTTP status code* and *description of status*. The description has **no** real effects and is kept as a convention as in HTTP, and it should not have any spaces. If the response has a content body, a Content-Length header should similarly be included, if there's content body. The corresponding responses for the above three requests could be:

- (a) 404_NotFound_
- (b) 200_OK_
- (c) 200_OK_content-length_6_ CS2105
- (d) 200_OK_content-length_1_0
- (e) 200_OK_
- (f) 200_OK_content-length_1_1

The server can assume that **all requests are correctly formatted** as defined above. However, as in a real web server, a request may contain multiple header fields, and some of them may be unknown to your server (but with valid format). Therefore, Content-Length and its value, if exist, may not be the first header field after the HTTP method and path, and it may not be the last as well. In addition, it is guaranteed that there are no duplicate header fields. Your server should parse all header fields and ignore irrelevant ones. Moreover, the maximum size of a content body is 5MB for this assignment.

The Key-value Store

A key-value store can be understood as a dictionary data structure that supports insertion, update, retrieval, and deletion of key-value pairs, with keys functioning as pointers to the values. For simplicity, we restrict keys to be case-sensitive ASCII strings in this assignment. Due to the use of `Content-Length` header, values can be any **binary string**. The server should keep all data in memory only and avoid accessing the disk.

The key-value store should be accessible to the client through paths with prefix `/key/`. To operate on a specific key, the complete HTTP path is `/key/` appended with the actual key string. For example, a request to operate on key `ModuleCode` specifies `/key/ModuleCode` as the HTTP path. The server should support the following key-value operations over the HTTP-like protocol:

1. Insertion and update

- (a) The HTTP request method should be `POST` and the value to be inserted (or updated) constitutes the content body. There should also be a `Content-Length` header indicating the number of bytes in the content body.
- (b) The server should respond with a `200 OK` status code after inserting or updating the value. The client always expects success statuses for insertions and updates.

2. Retrieval

- (a) The HTTP request method should be `GET`, and there is no content body.
- (b) If the key does not exist in the key-value store, the server returns a `404 NotFound` status code. Otherwise, the server should return a `200 OK` code, the correct `Content-Length` header and the value data in the content body.

3. Deletion

- (a) The HTTP request method is `DELETE`, and there is no content body.
- (b) If the key does not exist, the server returns a `404 NotFound` code. Otherwise, it should delete the key-value pair from the store and respond with a `200 OK` code and the deleted value string as the content body. The `Content-Length` header should also be sent accordingly.

The Counter Store

A counter store is very similar to key-value store, except that it stores non-negative frequency count, instead of bytes value, of a particular counter key. For simplicity, it should support update, retrieval of counter values, with counter keys functioning as pointers to the frequency values. Counter keys are case-sensitive ASCII strings. Unlike key-value store, no content-length header is present in a `POST` request, while it is needed in the response to a `GET` request. The server should keep all data in memory only and avoid accessing the disk.

The counter store should be accessible to the client through paths with prefix `/counter/`.

To operate on a specific counter key, the complete HTTP path is `/counter/` appended with the actual counter key string. For example, a request to operate on counter key `StudentNumber` specifies `/counter/StudentNumber` as the HTTP path. Note the keys in counter store is independent from the ones in key-value store, i.e. same key may appear in both stores and one stores bytes value while the other stores an integer. The server should support the following counter key operations over the HTTP-like protocol for a counter store:

1. Update

- (a) The HTTP request method should be `POST` without any content body. The frequency value of a corresponding counter key should be increased by 1. Any counter key not updated yet has frequency 0. Hence upon first update of a particular counter key, its frequency value becomes 1, and the next time 2, and so on. For simplicity, you may assume that the testing script will only request up to a maximum frequency of **9** for any key in the counter store.
- (b) The server should respond with a `200 OK` status code after updating the counter. The client always expects success status for updates.

2. Retrieval

- (a) The HTTP request method should be `GET`, and there is no content body.
- (b) If the counter key was never updated before, corresponding frequency should be 0. The server should return a `200 OK` code, the correct `Content-Length` header and the frequency value in the content body without leading 0s.

Grading Rubric

1. The server program is free from syntax errors and can be successfully executed (or compiled, for Java/C/C++). (1 mark)
2. The HTTP-like protocol (3 marks)
 - (a) The server parses valid HTTP requests and sends correctly formatted responses.
 - (b) The server supports persistent connection and correctly processes sequential requests as specified.
 - (c) The server works properly and responsively over intermittent connections simulated by the test script.
3. The key-value store (3 marks)
 - (a) The server supports insertion/update of key-value pairs through the `POST` method according to the spec.
 - (b) The server supports retrieval of value by key through the `GET` method according to the spec.
 - (c) The server supports removal of key-value pairs through the `DELETE` method according to the spec.

4. The counter store (1 mark)

- (a) The server supports update/retrieval of counter store according to the spec.

Common Errors/Issues

- The Key-Value store and the counter store are to be maintained throughout the execution of server and should not be reset for new connection.
- While sending the response avoid sending the response in multiple chunks. Create a single `bytearray()` before sending the response to client.
- Like in the case of Assignment 0, ensure that you handle header and the data separately
- Like in the case of Assignment 0, data is to be treated as raw bytes.
- Print your debug messages to `stdout`, we have redirected `stderr` to a temporary file, hence you will not see the printed message.
- There is a low possibility that the grading script is grading an existing running server program instead of yours, if they use same port. So start the assignment early to avoid "congestion" during last few days. An advice to ensure the grading feedback you received is on your current program is to print to screen your student number right after receiving a connection, and kill your own existing server programs if any before testing it again. The grading result will be based on your program if you see your student number printed for particular test case, or otherwise try again.